



# Intelligent Agents



**Instructor: Dr. Mohsin Ashraf**  
Slide credits: AI BERKELEY

# What is an Agent?

- “Intelligent agents continuously perform three functions: **perception** of dynamic conditions in the environment; **action** to affect conditions in the environment; and **reasoning** to interpret perceptions, solve problems, draw inferences, and determine actions.”



# Intelligent Agents

- An intelligent agent may learn from the environment to achieve their goals.
- Following are the main four rules for an AI agent:
  - **Rule 1:** An AI agent must have the ability to perceive the environment.
  - **Rule 2:** The observation must be used to make decisions.
  - **Rule 3:** Decision should result in an action.
  - **Rule 4:** The action taken by an AI agent must be a rational action.

# Example of Agents


- Human Agent
  - eyes, ears, skin, taste buds, etc. for sensors
  - hands, fingers, legs, mouth, etc. for actuators
- Robot
  - camera, infrared, bumper, etc. for sensors
  - grippers, wheels, lights, speakers, etc. for actuators
- Software Agent
  - functions as sensors
    - information provided as input to functions in the form of encoded bit strings or symbols
  - functions as actuators
    - results deliver the output

# Rational Agent

- **Rational Agent:**
  - For each possible percept sequence, a rational agent should select an action that is **expected to maximize its performance measure**, given the evidence provided by the percept sequence and whatever built-in knowledge the agent has.
- **Rationality** is distinct from **omniscience** (all-knowing with infinite knowledge)
  - **Actual (perfection) vs. Expected**
  - **Best vs. Optimal**

# Agent Function and Program

- **Behavior of Agent** – It is the action that agent performs after any given sequence of percepts
- Agent's behavior is mathematically described by
  - **Agent function**
    - A function mapping any given percept sequence to an action
- Practically it is described by
  - An **agent program**
  - The real implementation
- Problems:
  - what is “the right thing”
  - how do you measure the “best outcome”



# Specifying the Task Environment

Performance of Agents

# Performance of Agents

- Behaviour and performance of IA in terms of agent program.
  - Perception to Action Mapping
  - Ideal mapping: specifies which actions an agent ought to take at any point in time.
- Performance measure: a subjective measure to characterize how successful an agent is (e.g., speed, power, accuracy, etc.)
- **PEAS:**
  - **Performance measure,**
  - **Environment,**
  - **Actuators,**
  - **Sensors**




# PEAS: Automated Taxi

- Performance measure
  - Correct Destination, Fuel, time, cost, violations, safety, comfort, profit
- Environment
  - rural, urban, other drivers, customers
- Actuators
  - Steering, brake, fuel, display/speaker
- Sensors
  - Camera, radar, accelerometer, engine sensors, microphone



# PEAS: Medical diagnosis system

- Performance measure:
  - Healthy patient, minimize costs, lawsuits
- Environment:
  - Patient, hospital, staff
- Actuators:
  - Screen display (questions, tests, diagnoses, treatments, referrals)
- Sensors:
  - Keyboard (entry of symptoms, findings, patient's answers)



# PEAS: Part-picking robot

- Performance measure:
  - Percentage of parts in correct bins
- Environment:
  - Conveyor belt with parts, bins
- Actuators:
  - Jointed arm and hand
- Sensors:
  - Camera, joint angle sensors




# PEAS: Interactive English tutor

- Performance measure:
  - Maximize student's score on test
- Environment:
  - Set of students
- Actuators:
  - Screen display (exercises, suggestions, corrections)
- Sensors:
  - Keyboard

# Environment Types

# Properties of environments




- **Fully observable vs. partially observable**
- **Or Accessible vs. inaccessible**
  - Can the agent observe the complete state of the environment?
  - If an agent's sensors gives it access to the complete state of the environment, then we say that environment is fully observable to the agent.
  - A fully observable environment is convenient because the agent need not maintain any internal state to keep track of the world.
  - *Is the current state of the world fully known at each step?*

# Properties of environments



- **Deterministic vs. nondeterministic (Stochastic).**
  - Is there uncertainty in how the world works?
  - If the next state of the environment is completely determined by the current state and the actions selected by the agents, then we say the environment is deterministic.
  - If the environment is inaccessible, then it may *appear* to be nondeterministic (bunch of uncertainties).
- Taxi driving is clearly stochastic in this sense, because one can never predict the behavior of traffic exactly.

# Environment Types



vs.



- **Episodic vs. Sequential:**

- **Episodic: next episode doesn't depend on previous actions.**
- The agent's experience is divided into atomic "episodes"
- each episode consists of the agent perceiving and then performing a single action, and
- the choice of action in each episode depends only on the episode itself.
- the next episode does not depend on the actions taken in previous episodes
- For example, an agent that has to spot defective parts on an assembly line (EPISODIC)
- Chess and Taxi driver (SEQUENTIAL).



# Environment Types

- **Static vs. Dynamic:**

- If the environment can change while an agent is performing action, then we say the environment is dynamic for that agent; otherwise, it is static.
- Dynamic environments are continuously asking the agent what it wants to do
- Static environments are easy to deal with, because the agent does not keep on looking at the environment while it is deciding on an action.
- The environment is **semi-dynamic** if the environment itself does not change with the passage of time but the agent's performance score does.

# Environment Types

- **Discrete vs. continuous:**

- Is there a finite (or countable) number of possible environment states?
- The environment is discrete if the number of actions and possible states of the environment is finite otherwise it is continuous.
- If there are a limited number of distinct, clearly defined percepts and actions, we say that the environment is discrete.
  - Chess, since there are a fixed number of possible moves on each turn.
  - Taxi driving is continuous.

# Types of Environment

- **Single Agent vs. Multi Agent**

- Is the agent the only thing acting in the world?
- In the single agent environment there is only one agent
  - A computer software playing crossword puzzle
- In multiagent systems, there are more than one active agents
  - Video games



VS.






# Environment Types Example

Crossword puzzle



EclipseCrossword.com



**Crossword puzzle**

**Chess**

**Taxi driving**

**Part-picking robot**

**English tutor**

**Observable**

**Yes**

**Yes**

**No**

**No**

**No**

**Deterministic**

**Yes**

**Yes**

**No**

**No**

**No**

**Episodic**

**No**

**No**

**No**

**Yes**

**No**

**Static**

**Yes**

**Yes**

**No**

**No**

**No**

**Discrete**

**Yes**

**Yes**

**No**

**No**

**Yes**

**Agents**

**Single**

**Multi**

**Multi**

**Single**

**Multi**

# Structure of Agents

# Agent design

- **The environment type largely determines the agent design**
  - *Fully/partially observable* => agent requires **memory** (internal state)
  - *Discrete/continuous* => agent may not be able to enumerate **all states**
  - *Stochastic/deterministic* => agent may have to prepare for **contingencies**
  - *Single-agent/multi-agent* => agent may need to behave **randomly**

# Skeleton Agent Program

- Basic framework for an agent program

```
function SKELETON-AGENT(percept) returns action
```

```
  static: memory
```

```
    memory      := UPDATE-MEMORY(memory, percept)
```

```
    action      := CHOOSE-BEST-ACTION(memory)
```

```
    memory      := UPDATE-MEMORY(memory, action)
```

```
return action
```

# Table Agent Program

- agent program based on table lookup

```
function TABLE-DRIVEN-AGENT(percept) returns action  
  static: percepts // initially empty sequence*  
           table    // indexed by percept sequences  
                // initially fully specified  
  
  append percept to the end of percepts  
  action      := LOOKUP(percepts, table)  
  
return action
```

\* Note: the storage of percepts requires writeable memory

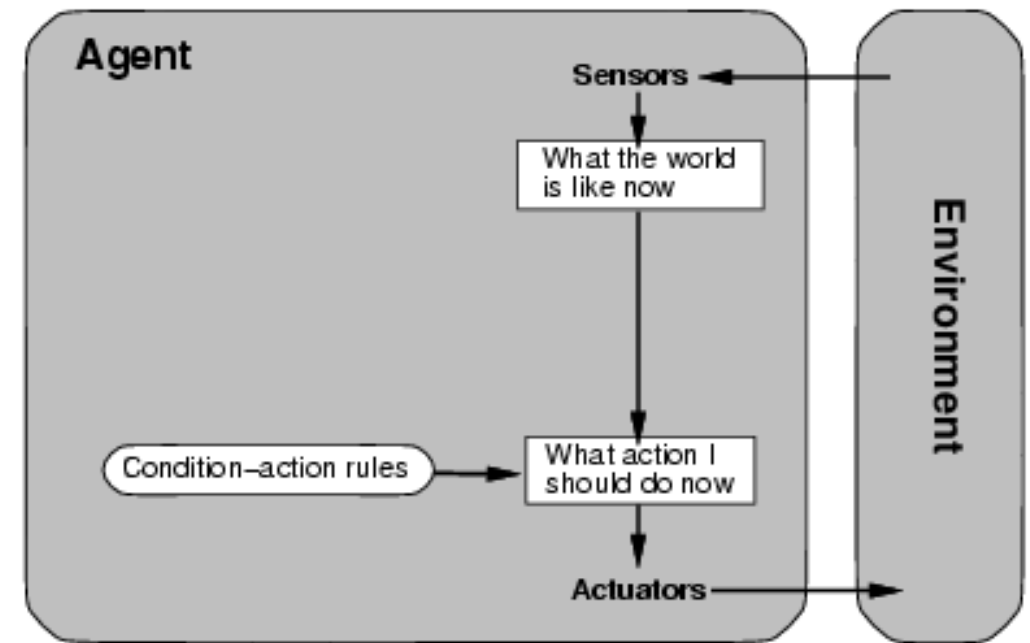


# Structure of Agent

- Different ways of achieving the mapping from percepts to actions
- Types of agents (increasing in generality and ability to handle complex environments)
  - Simple reflex agents
  - Model-based reflex agents
    - keep track of the world
  - Goal-based agents
    - work towards a goal
  - Utility-based agents
  - Learning agent

# Simple reflex agents

- Act only on the basis of the **current percept, ignoring the rest of the percept history**.
- The agent function is based on the ***condition-action rule***: if condition then action.
- the vacuum agent is a simple reflex agent, because its decision is based only on the current location and on whether that location contains dirt.
- Fully-observable




# Simple Reflex Agents

- The Simple reflex agents are the simplest agents. These agents take decisions on the basis of the current percepts and ignore the rest of the percept history.
- These agents only succeed in the fully observable environment.
- The Simple reflex agent works on Condition-action rule, which means it maps the current state to action.
- Problems for the simple reflex agent design approach:
  - They have very limited intelligence
  - Mostly too big to generate and to store.
  - Not adaptive to changes in the environment.

# Example: Vacuum Agent

- **Performance?**
  - 1 point for each square cleaned in time T?
  - #clean squares per time step - #moves per time step?
- **Environment:** vacuum, dirt, multiple areas defined by square regions
- **Actions:** left, right, suck, idle
- **Sensors:** location and contents
  - [A, dirty]
- Rational is not omniscient
  - Environment may be partially observable
- Rational is not clairvoyant
  - Environment may be stochastic
- Thus Rational is not always successful



If status=Dirty then return Suck  
else if location=A then return Right  
else if location=B then right Left

# Simple Reflex Agent

```
function SIMPLE-REFLEX-AGENT(percept) returns an action
  persistent: rules, a set of condition–action rules

  state  $\leftarrow$  INTERPRET-INPUT(percept)
  rule  $\leftarrow$  RULE-MATCH(state, rules)
  action  $\leftarrow$  rule.ACTION
  return action
```

**Figure 2.10** A simple reflex agent. It acts according to a rule whose condition matches the current state, as defined by the percept.

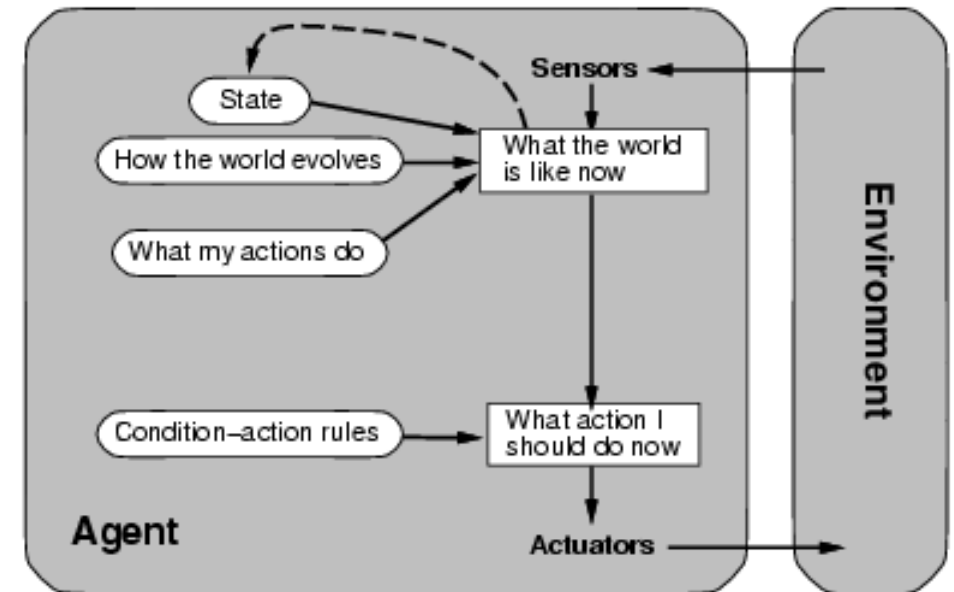
```
1  def simple_reflex_agent(percept):
2      state = get_state_from_percept(percept)
3      rule = match_rule(state, rules)
4      action = rule.Action
5      return action
```

simple\_reflex\_agent.py hosted with ❤ by GitHub

[view raw](#)

# Model-based reflex agents

- Keep track of the part of the worlds it can't see now.
- Its current state is stored inside the agent ---- **Internal state**
- Percept history and impact of action on the environment can be determined by using internal model.
- It then chooses an action in the same way as reflex agent.
- Partially-observable



# Model-Based Reflex Agents

- The Model-based agent can work in a partially observable environment, and track the situation.
- A model-based agent has two important factors:
  - **Model:** It is knowledge about "how things happen in the world," so it is called a Model-based agent.
  - **Internal State:** It is a representation of the current state based on percept history.
- These agents have the model, "which is knowledge of the world" and based on the model they perform actions.
- Updating the agent state requires information about:
  - How the world evolves
  - How the agent's action affects the world.

# Model-Based Agent

**function** MODEL-BASED-REFLEX-AGENT(*percept*) **returns** an action  
**persistent:** *state*, the agent's current conception of the world state  
              *model*, a description of how the next state depends on current state and action  
              *rules*, a set of condition–action rules  
              *action*, the most recent action, initially none

*state*  $\leftarrow$  UPDATE-STATE(*state*, *action*, *percept*, *model*)  
*rule*  $\leftarrow$  RULE-MATCH(*state*, *rules*)  
*action*  $\leftarrow$  *rule*.ACTION  
**return** *action*

**Figure 2.12** A model-based reflex agent. It keeps track of the current state of the world, using an internal model. It then chooses an action in the same way as the reflex agent.

```
1 def model_based_reflex_agent(percept):
2     state = update_state(state, action, percept, model)
3     rule = match_rule(state, rules)
4     action = rule.Action
5     return action
```

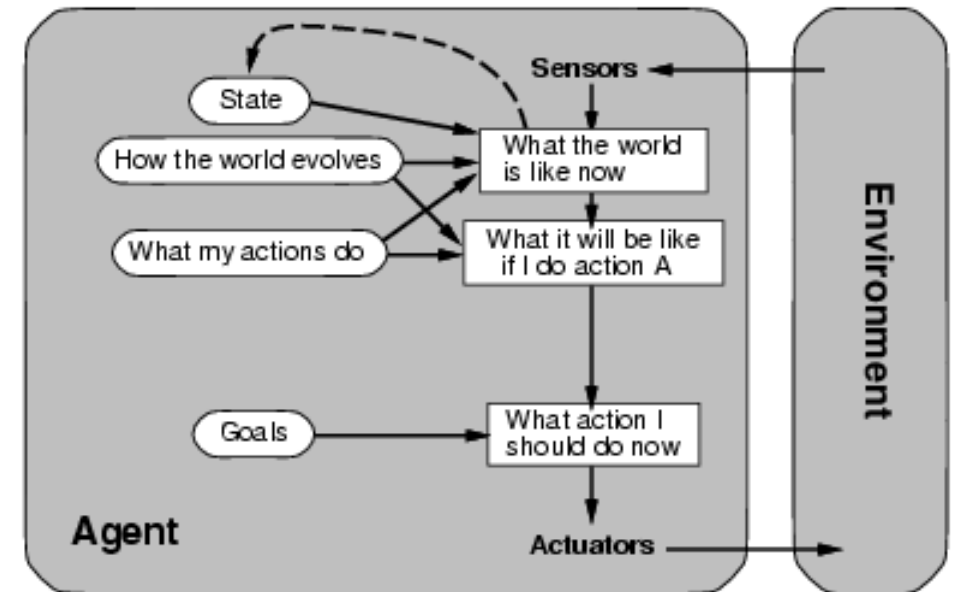
model\_based\_reflex\_agent.py hosted with ❤ by GitHub

[view raw](#)



# Goal-based agents

- Expand on the capabilities of the model-based agents, by using "goal" information.
- Goal information describes situations that are desirable.
- This allows the agent a way to choose among multiple possibilities, selecting the one which reaches a goal state.
- Search and planning



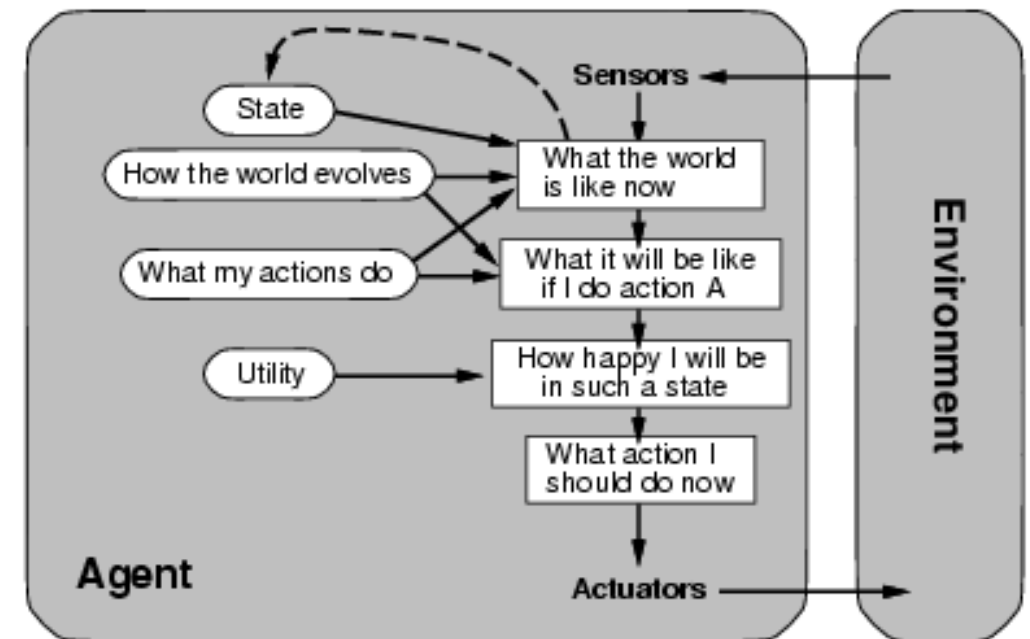
# Goal Based Agents

- The knowledge of the current state environment is not always sufficient to decide for an agent to what to do.
- The agent needs to know its goal which describes desirable situations.
- Goal-based agents expand the capabilities of the model-based agent by having the "goal" information.
- They choose an action, so that they can achieve the goal.
- These agents may have to consider a long sequence of possible actions before deciding whether the goal is achieved or not. Such considerations of different scenario are called searching and planning, which makes an agent proactive.

```
function Goal-Based-Agent(percept, goal) returns action  
  static : rules, a set of condition-action rules  
           state, a description of the current state  
  state := Update-State (state, percept)  
  rule := Plan-Best-Move(state, rules, goal)  
  action := Rule-Action[rule]  
  state := Update-State (state, action)  
return action
```

# Utility-based agents

- We say that if one world state is preferred to another, then it has higher *utility* for the agent.
- Utility is a function that maps a state onto a real number.
  - $\text{state} \rightarrow \mathbb{R}$
- Any rational agent possesses a utility function.
- A utility function maps each state after each action to a real number representing how efficiently each action achieves the goal.

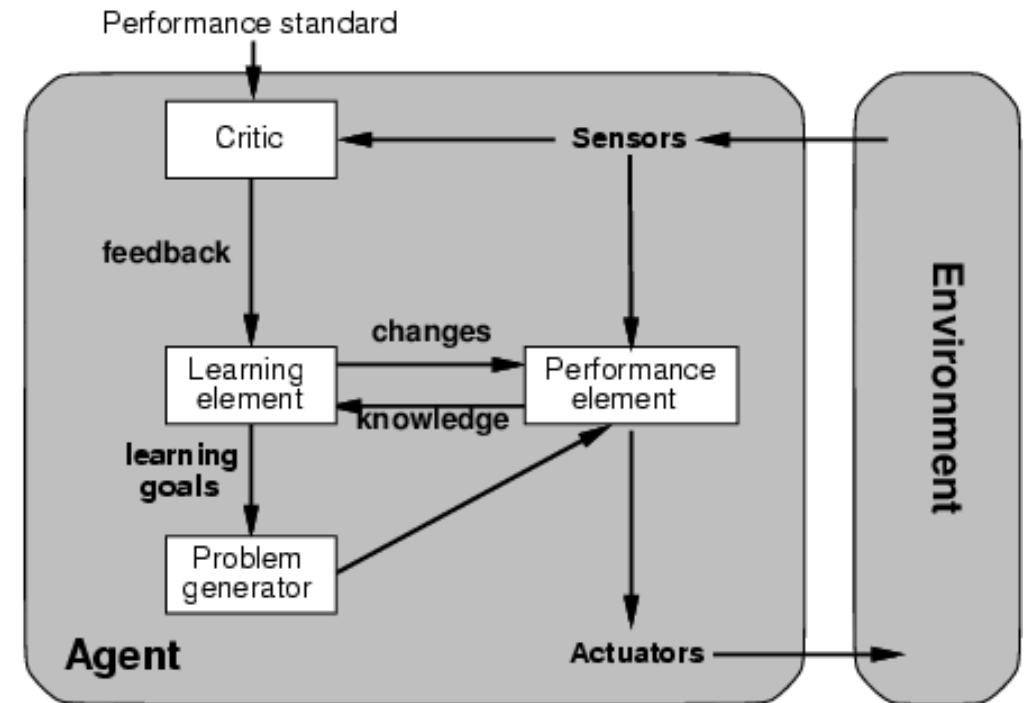


# Utility Based Agents

- These agents are similar to the goal-based agent but provide an extra component of utility measurement which makes them different by providing a measure of success at a given state.
- Utility-based agent act based not **only goals but also the best way to achieve the goal.**
- The Utility-based agent is useful when there are multiple possible alternatives, and an agent has to choose in order to perform the best action.
- The utility function maps each state to a real number to check how efficiently each action achieves the goals.

# Learning Agents

- It allows agents to operate in unknown environments and to become more competent.
- learning element
  - responsible for making improvements
- performance element
  - responsible for selecting actions.
- Uses feedback from the **critic** on how the agent is doing.
- **problem generator**
  - responsible for suggesting actions that will lead to new experiences suggest (exploratory actions).



# Learning Agents

- A learning agent in AI is the type of agent which can learn from its past experiences, or it has learning capabilities.
- It starts to act with basic knowledge and then able to act and adapt automatically through learning.
- A learning agent has mainly four conceptual components, which are:
  - **Learning element:** It is responsible for making improvements by learning from environment
  - **Critic:** Learning element takes feedback from critic which describes that how well the agent is doing with respect to a fixed performance standard.
  - **Performance element:** It is responsible for selecting external action
  - **Problem generator:** This component is responsible for suggesting actions that will lead to new and informative experiences.
- Hence, learning agents are able to learn, analyze performance, and look for new ways to improve the performance.

# Intelligent Agents

- **Simple reflex agents** respond directly to percepts
- **model-based reflex agents** maintain internal state to track aspects of the world that are not evident in the current percept.
- **Goal-based agents** act to achieve their goals,
- **utility-based agents** try to maximize their own expected “happiness.”