

Informed Search - II

A* Search Algorithm



Previous Lecture

- Heuristics
- Greedy Search

eval-fn: $f(n) = g(n) + h(n)$

A* Search


eval-fn: $f(n) = g(n) + h(n)$

A* (A Star)

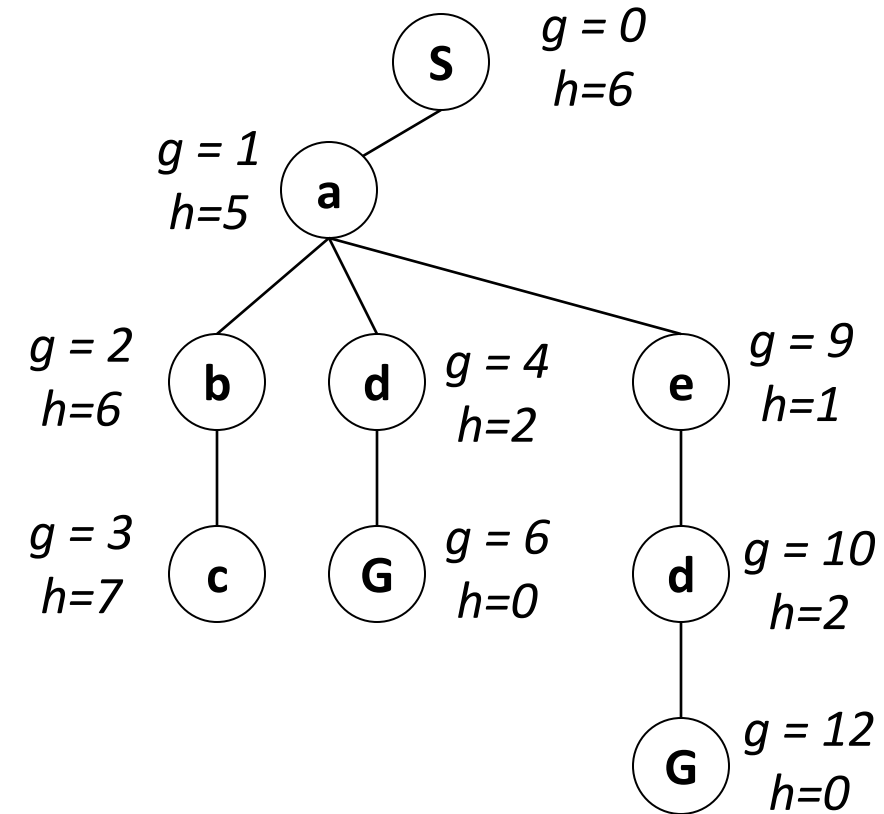
- Although Greedy Search can considerably cut the search time (**efficient**), it is **neither optimal nor complete**.
- Uniform Cost Search minimizes the cost $g(n)$ from the initial state to n . **UCS is optimal and complete but not efficient**.
- New Strategy: **Combine Greedy Search and UCS** to get an **efficient algorithm** which is **complete** and **optimal**.
- A* uses a heuristic function which combines $g(n)$ and $h(n)$:
$$f(n) = g(n) + h(n)$$

Combining UCS and Greedy


- Uniform-cost orders by path cost, or *backward cost* $g(n)$
- Greedy orders by goal proximity, or *forward cost* $h(n)$



- **A* Search** orders by the sum: $f(n) = g(n) + h(n)$



A* Search



State	Heuristic: h(n)
A	366
B	374
C	329
D	244
E	253
F	178
G	193
H	98
I	0


$$f(n) = g(n) + h(n)$$

g(n): is the exact cost to reach node n from the initial state.

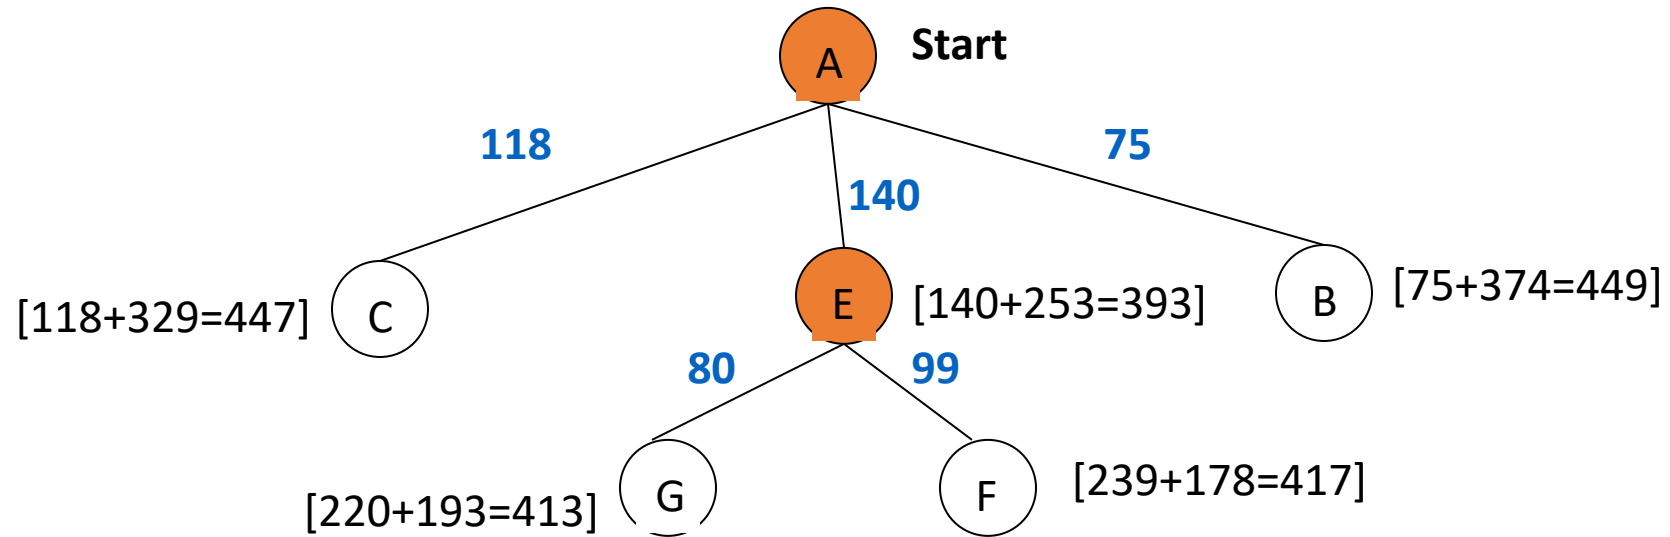
A* Search: Tree Search



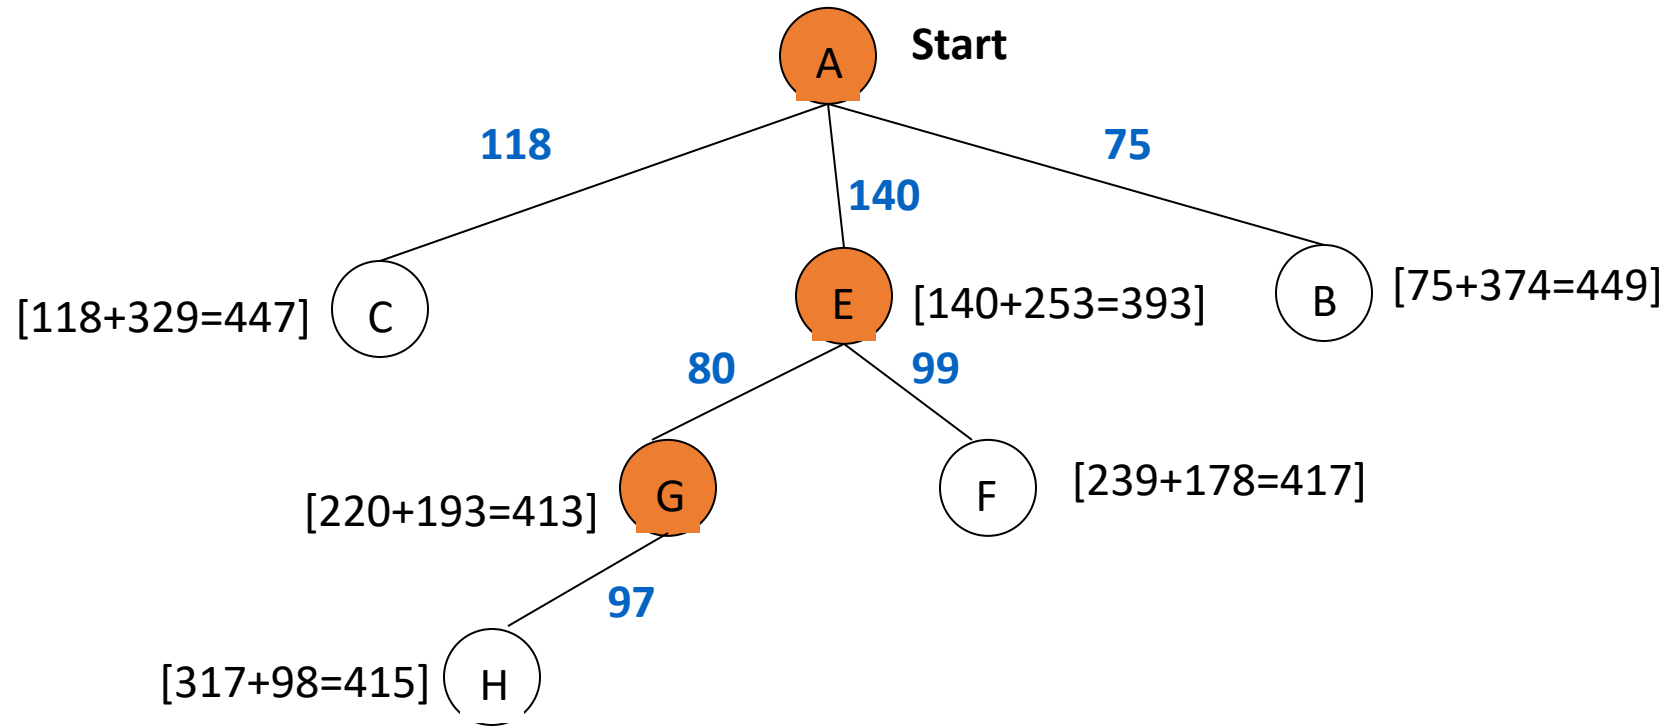
A* Search: Tree Search



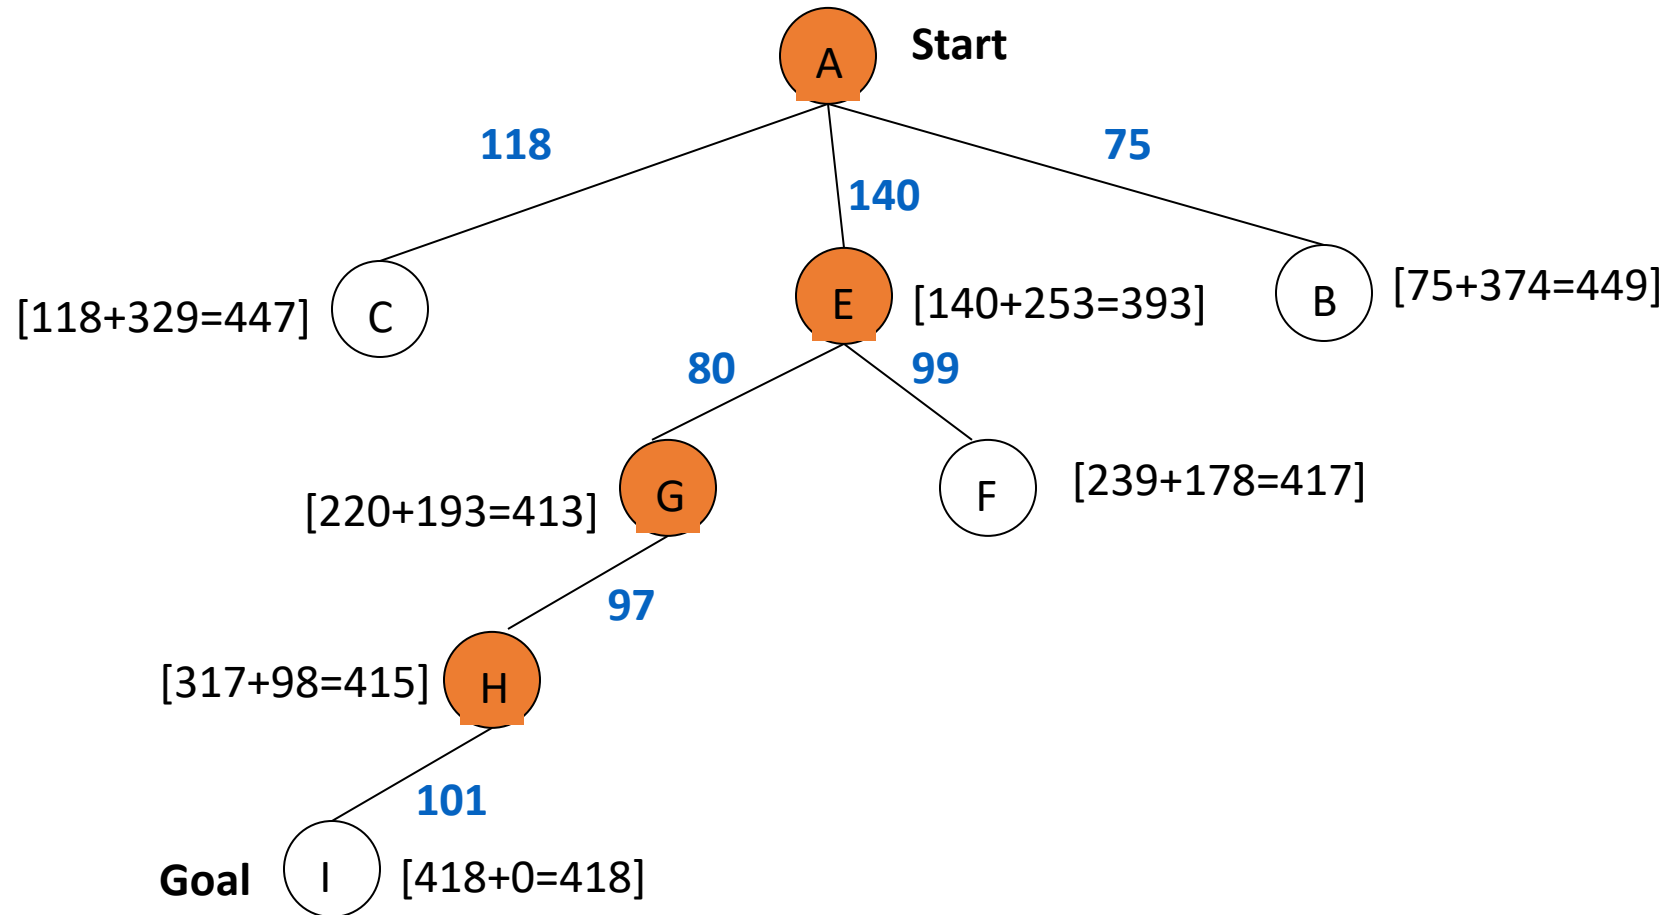
A* Search: Tree Search



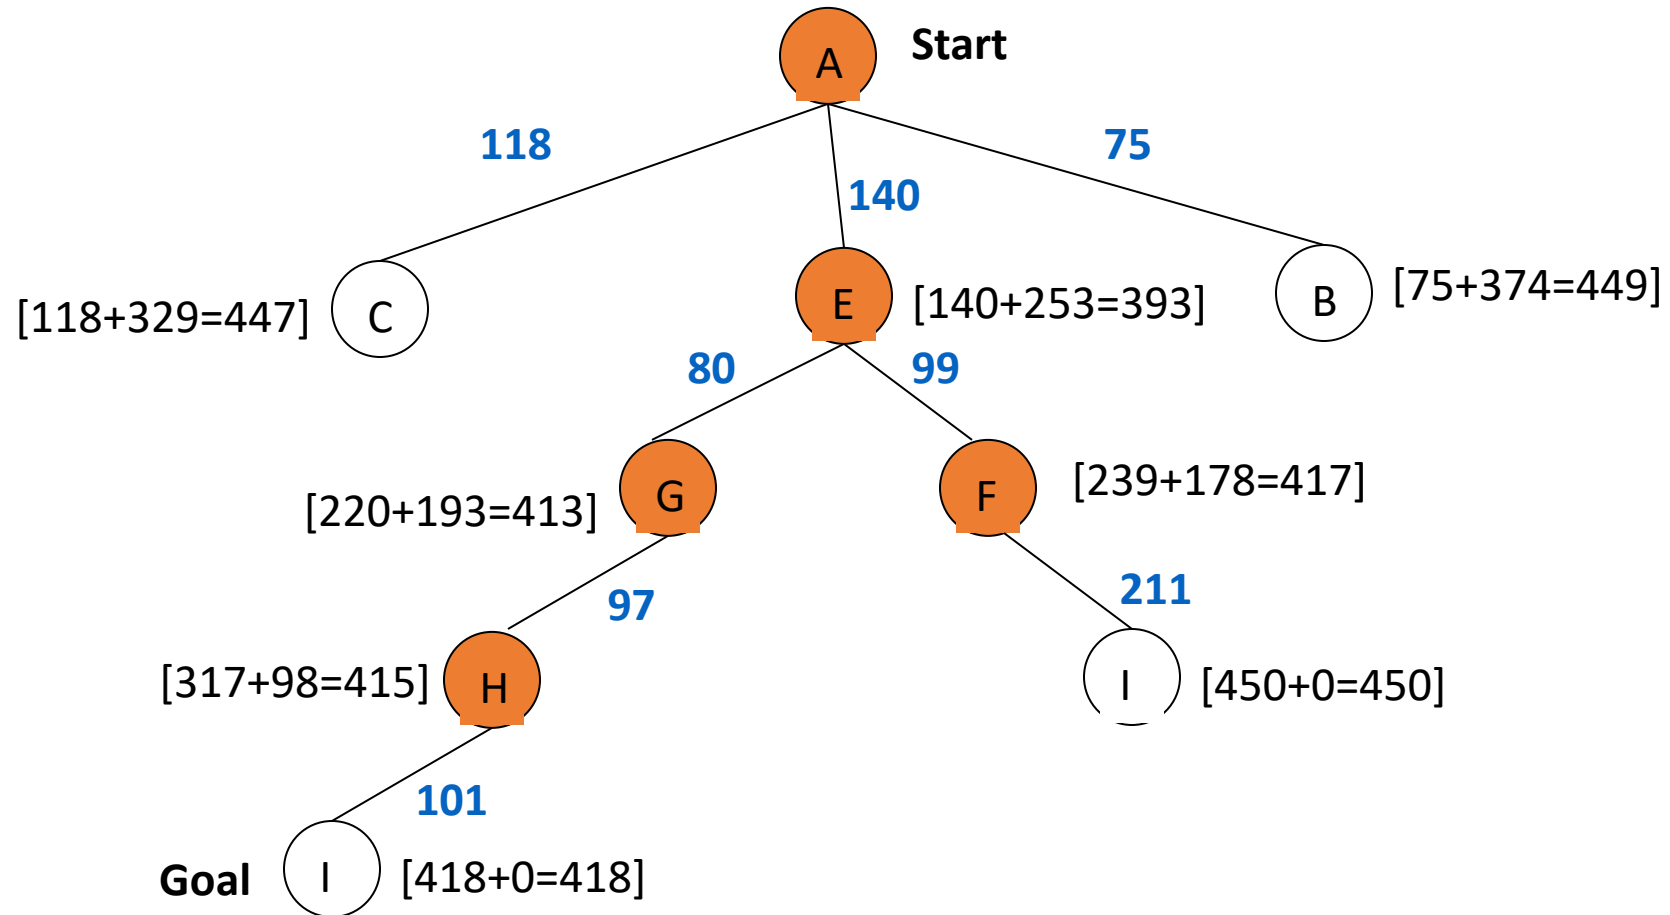
A* Search: Tree Search



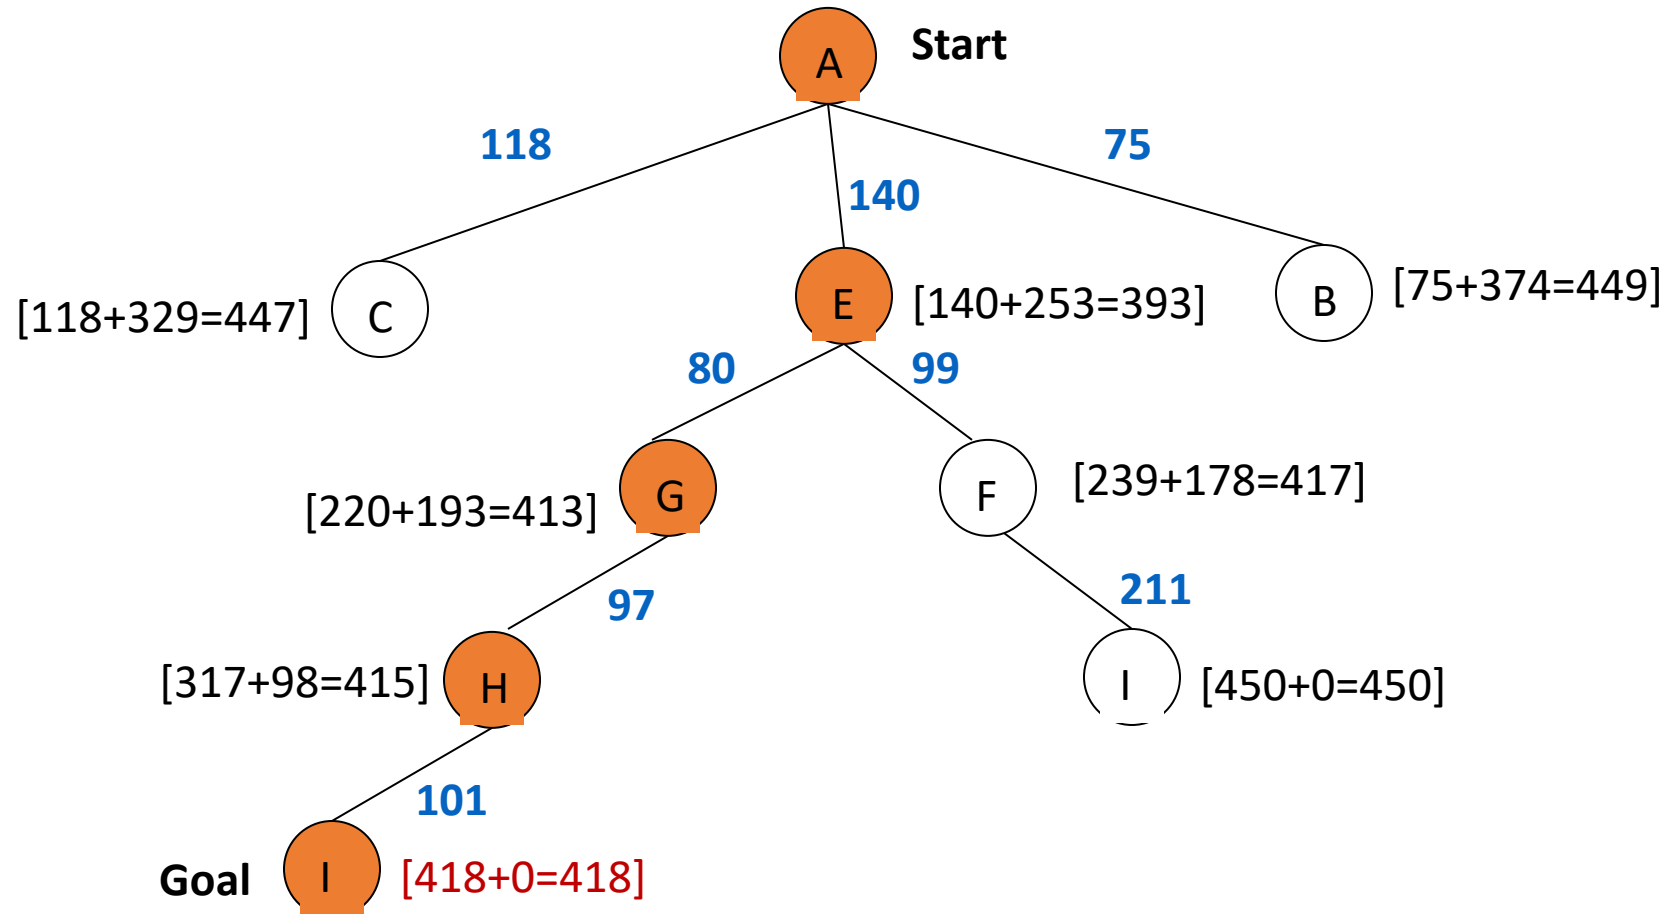
A* Search: Tree Search



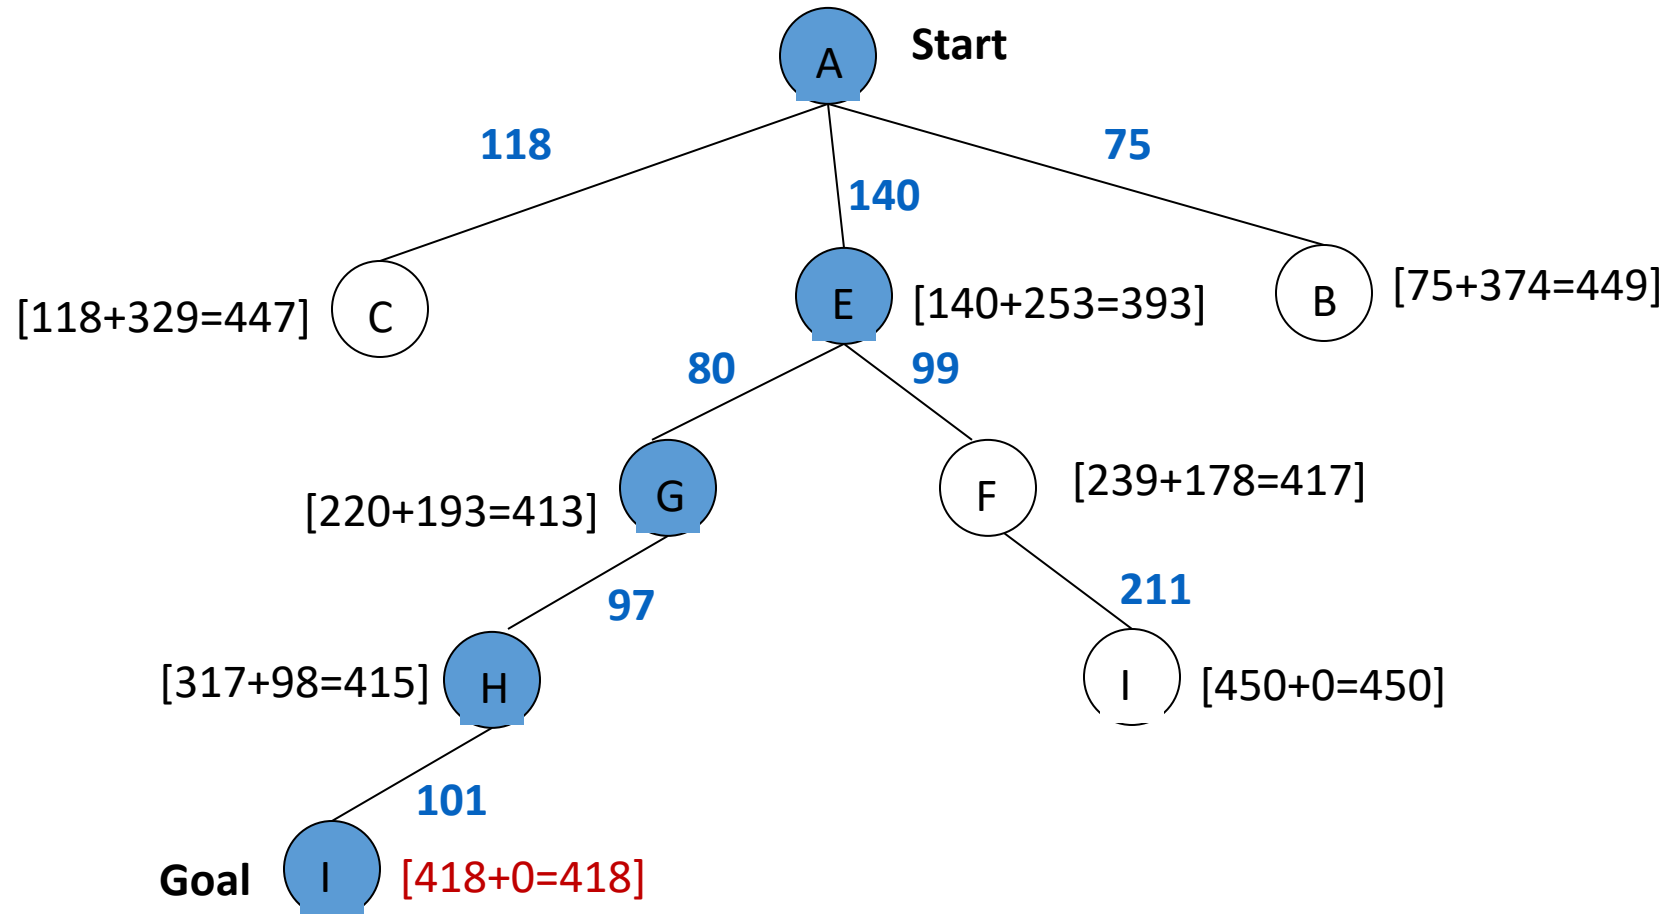
A* Search: Tree Search



A* Search: Tree Search




A* Search: Tree Search




$$\text{dist}(A-E-G-H-I) = 140 + 80 + 97 + 101 = 418$$

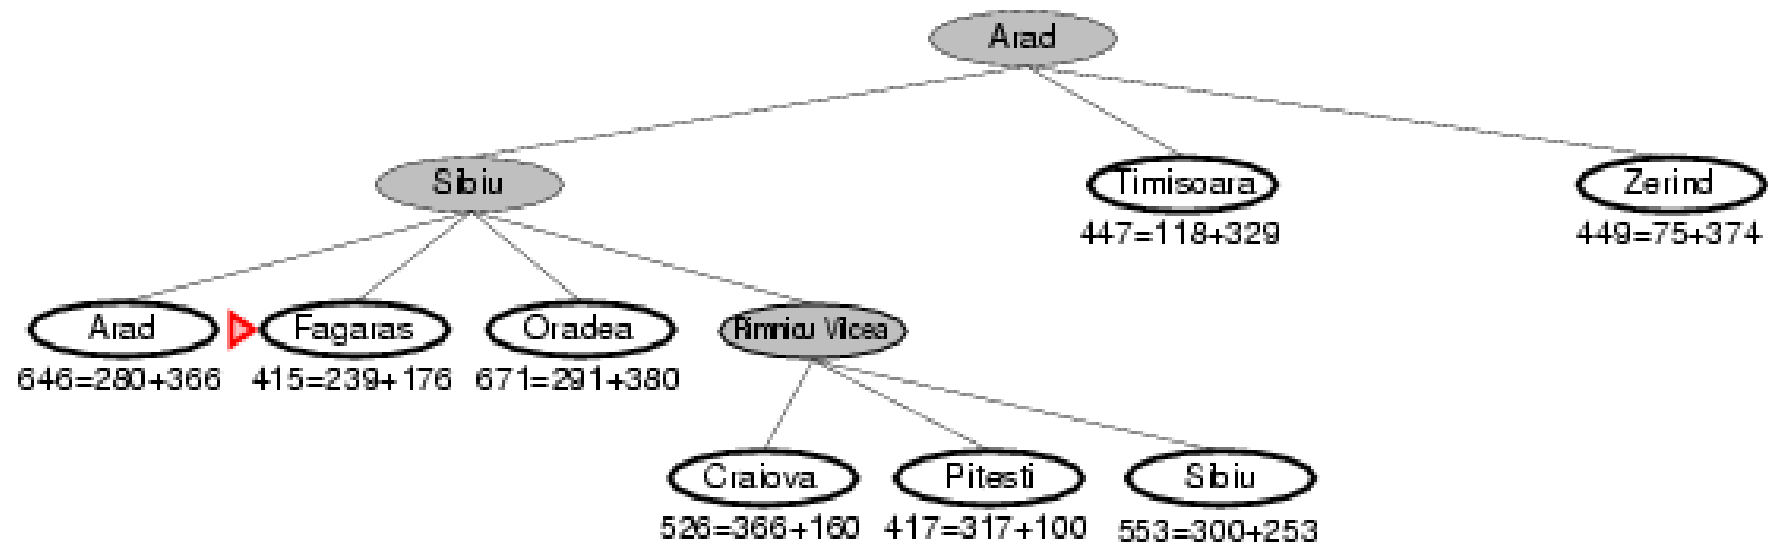
A* search example



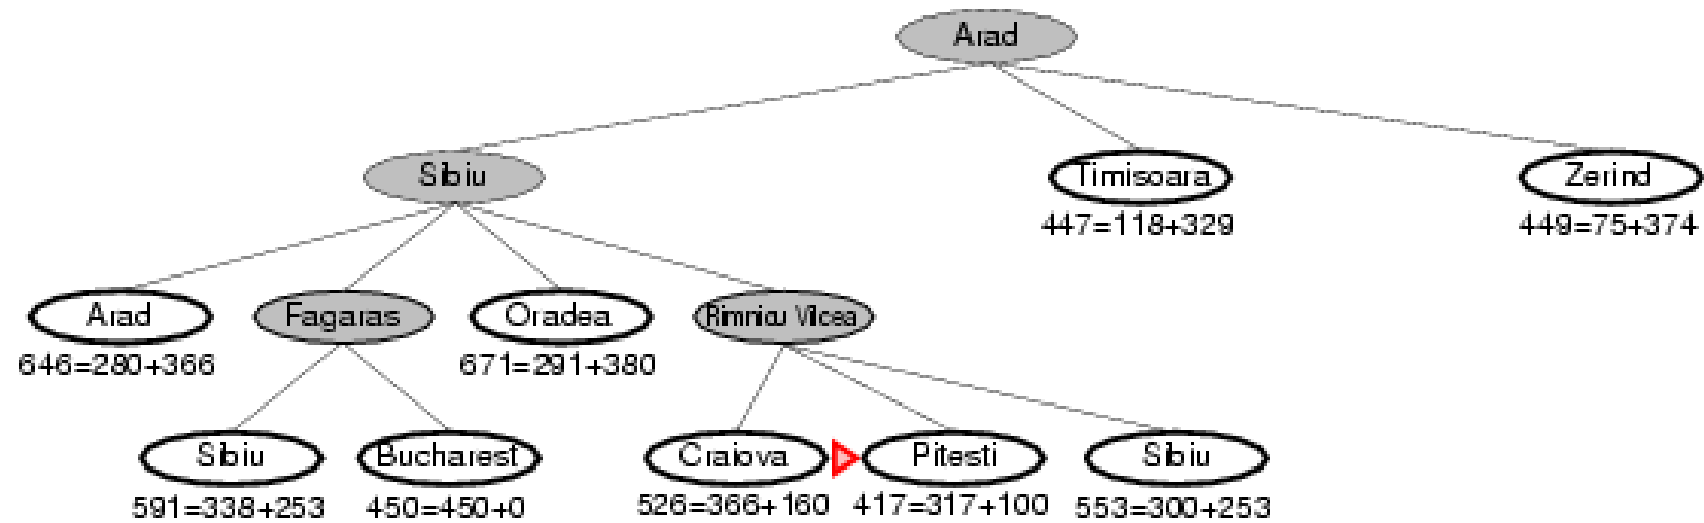
A* search example



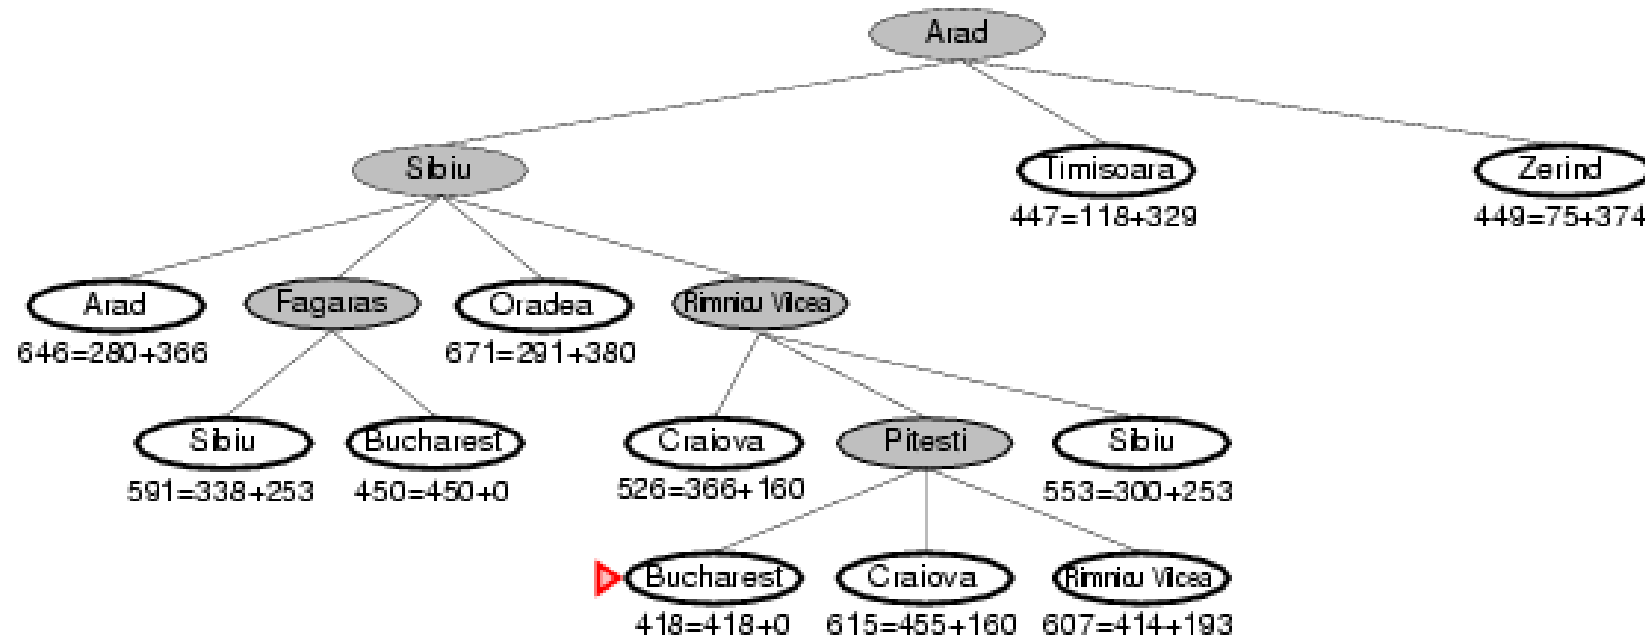
A* search example




A* search example



A* search example




Heuristic Value	
S	7
A	6
B	2
C	1
D	0




$f(n) = g(n) + h(n)$

	FRONTIER	EXPAND	EXPLORED
1	(S,0+7=7)	S	Empty
2	(S-A,1+6=7)(S-B,4+2=6)	B	S
3	(S-A,7)(S-B-C,6+1=7)	A	S,B
4	(S-A-B,3+2=5) (S-A-C,6+1=7)(S-A-D,13+0=13)(S-B-C,7)	C	S,B,A
5	(S-A-C-D,9+0=9)	D	S,B,A,C
6	OR (S-B-C-D,9+0=9)		
7			



Is A* Optimal?



- What went wrong?
 - Actual bad goal cost < estimated good goal cost
 - We need estimates to be less than actual costs!

Admissible Heuristic Functions

- A heuristic $h(n)$ is admissible if $h(n)$ never overestimates the cost to reach the goal.
 - Admissible heuristics are “optimistic” approximation

- Mathematically h is admissible if

$$\text{For all } n, h(n) \leq C(n)$$

where

- n is the node, h is heuristic
- $h(n)$ is the cost indicated by h to reach goal from n
- $C(n)$ is actual cost to reach a goal from n


Admissible heuristics Examples

- $h_{SLD}(n)$ Straight line distance heuristic
 - Because never overestimate the actual road distance.
- Number of misplaced tiles in n-puzzle problem

Admissible Heuristic Functions

- Evaluation function $f(n) = g(n) + h(n)$ is **admissible** if $h(n)$ is admissible.
- However, $g(n)$ is the exact cost to reach node n from the initial state.
- Therefore, $f(n)$ never over-estimate the true cost to reach the goal state through node n .
- **Theorem:** If $h(n)$ is admissible, A^* using TREE-SEARCH is optimal

A* Search: if h not admissible



State	Heuristic: h(n)
A	366
B	374
C	329
D	244
E	253
F	178
G	193
H	138
I	0


Goal $f(n) = g(n) + h(n) - \text{(H-I) Overestimated}$

$g(n)$: is the exact cost to reach node n from the initial state.

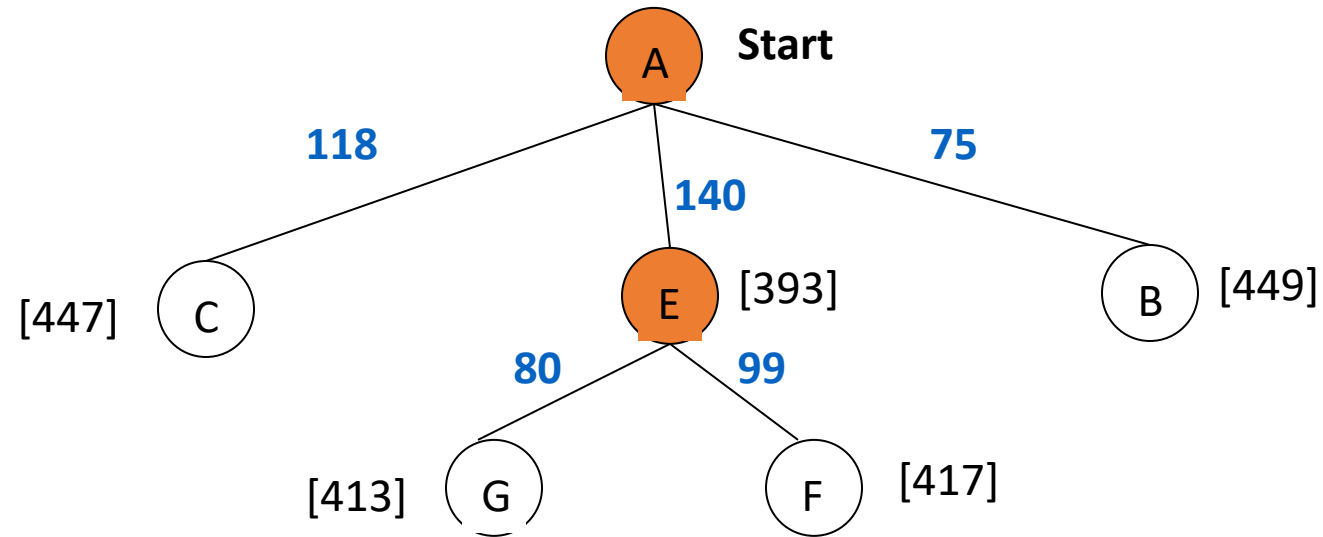
A* Search: Tree Search



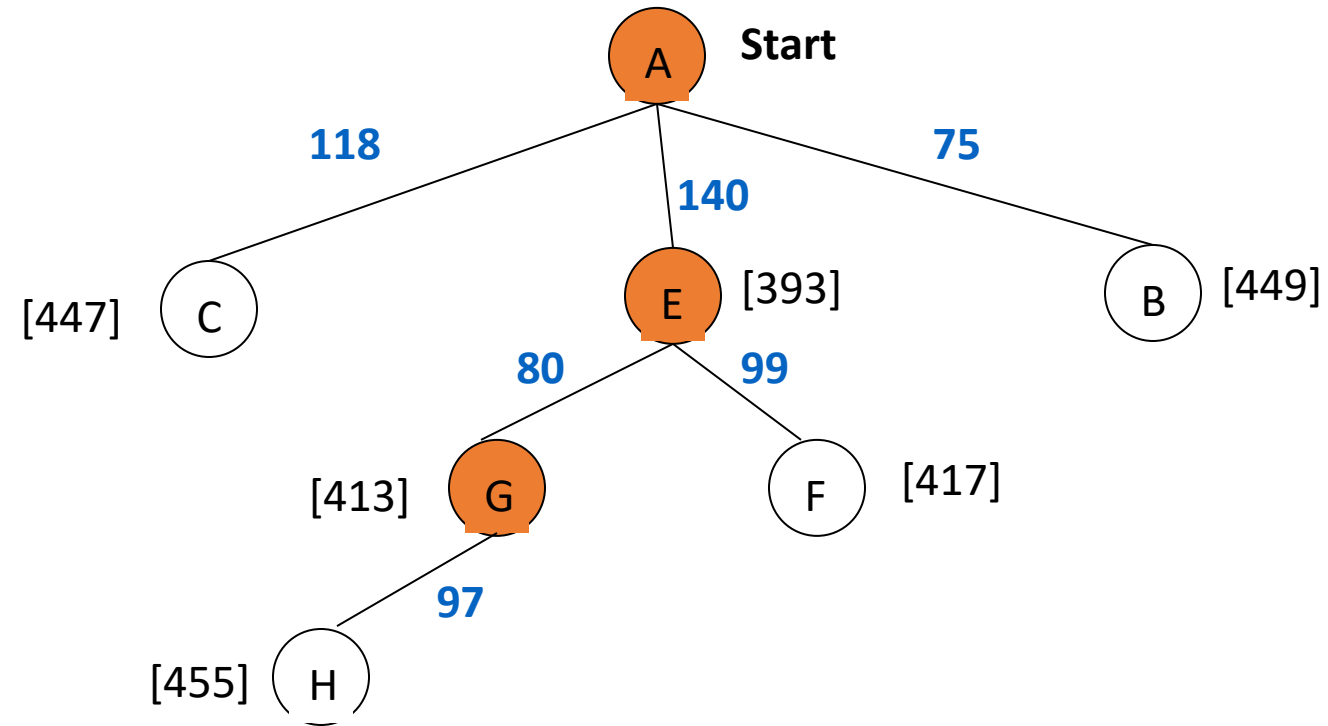
A* Search: Tree Search



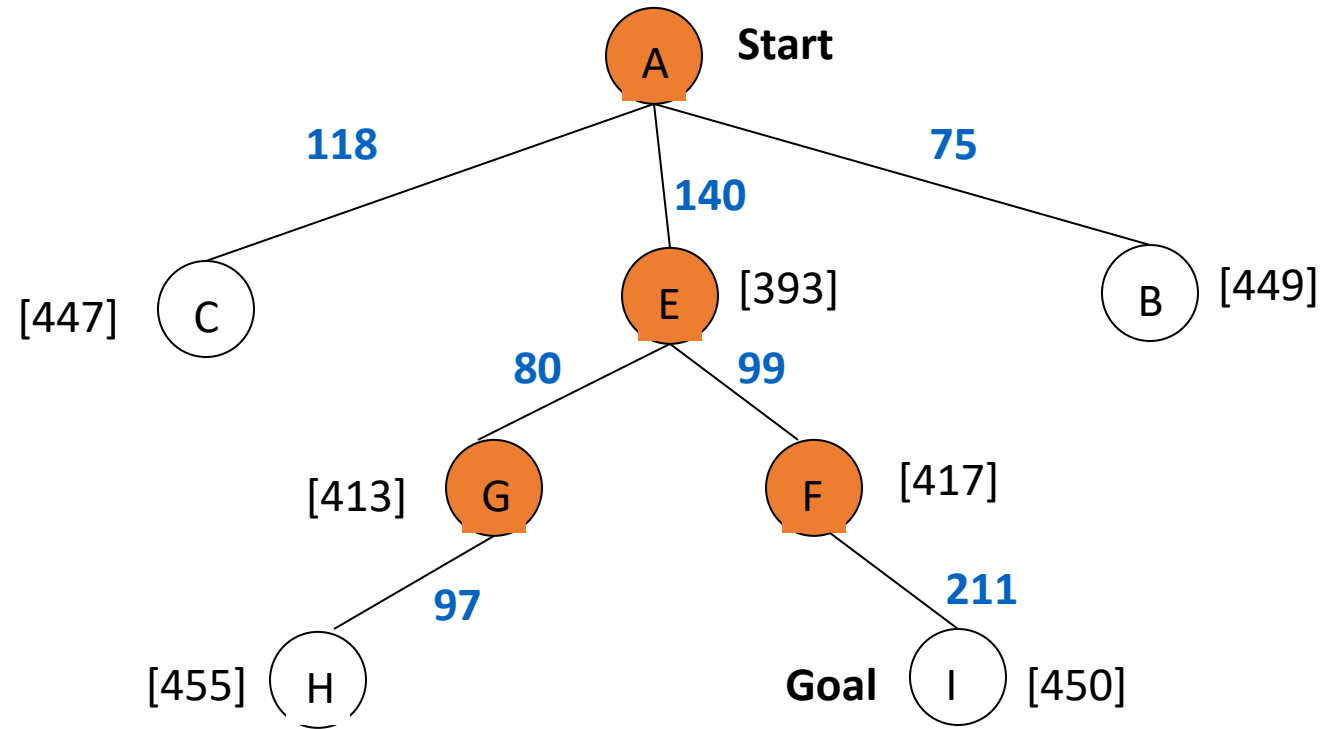
A* Search: Tree Search



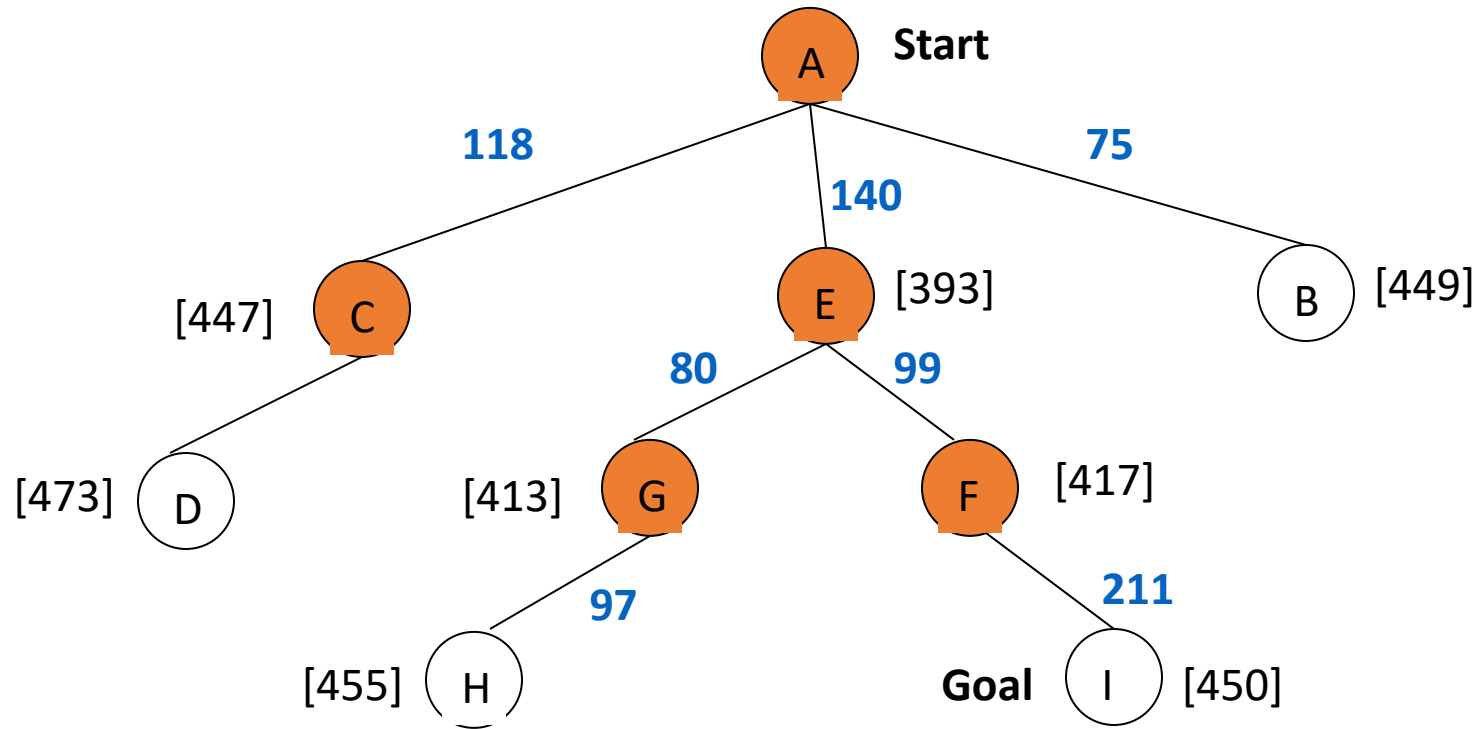
A* Search: Tree Search



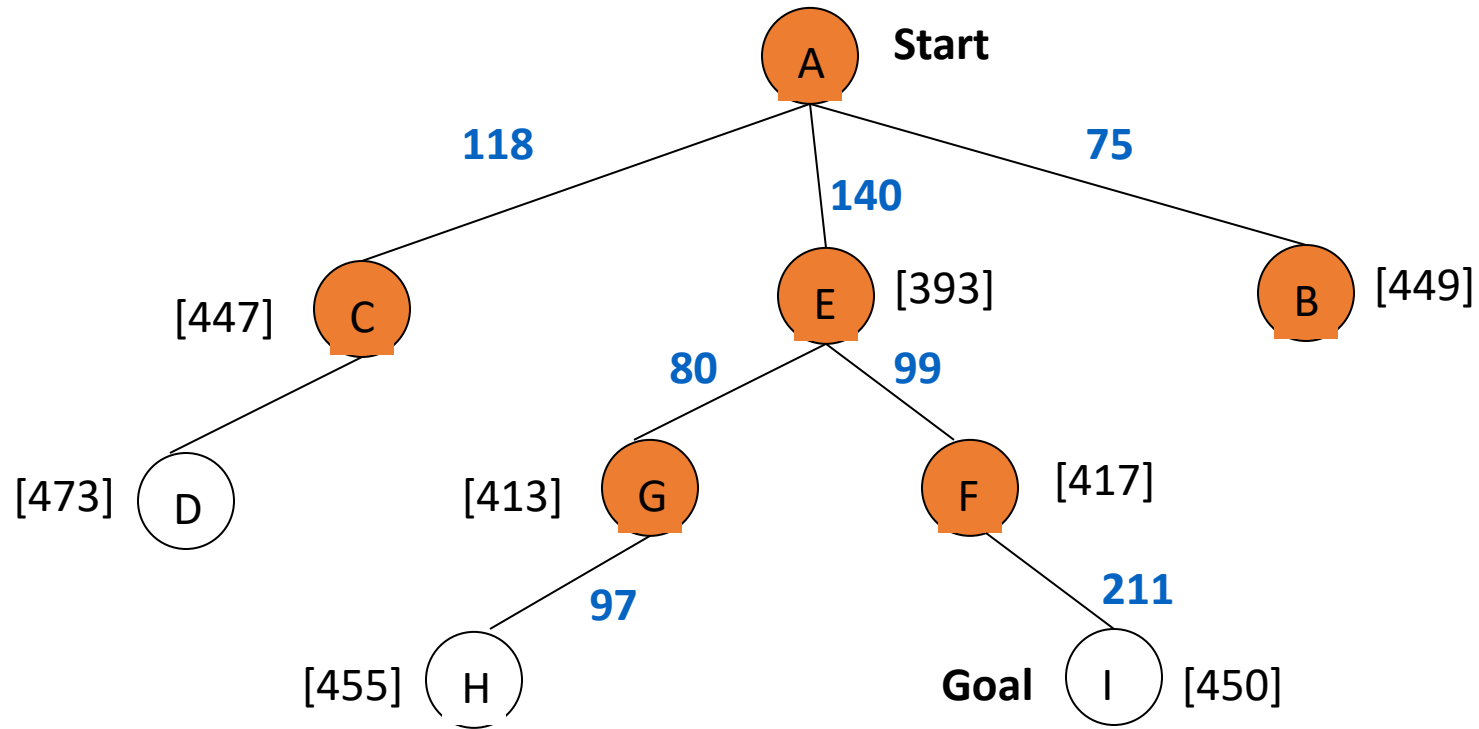
A* Search: Tree Search



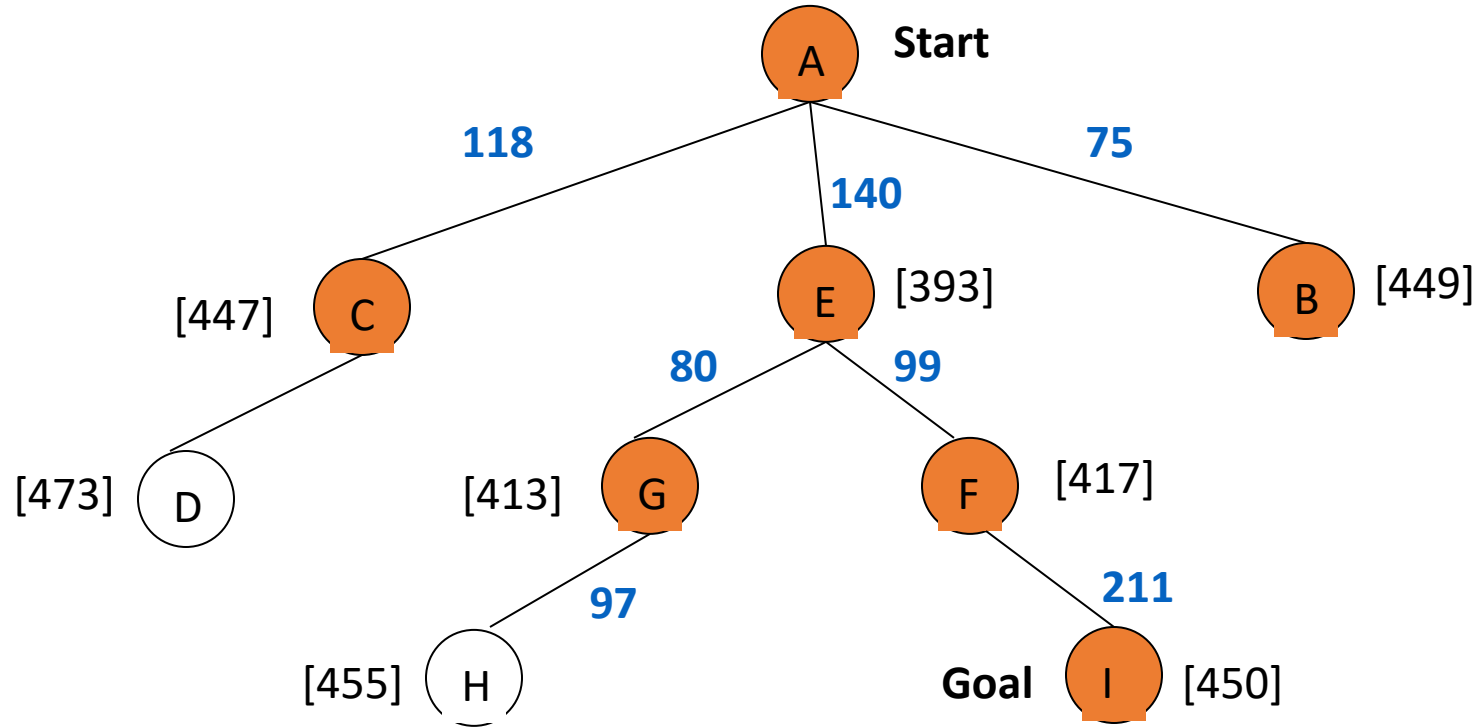
A* Search: Tree Search



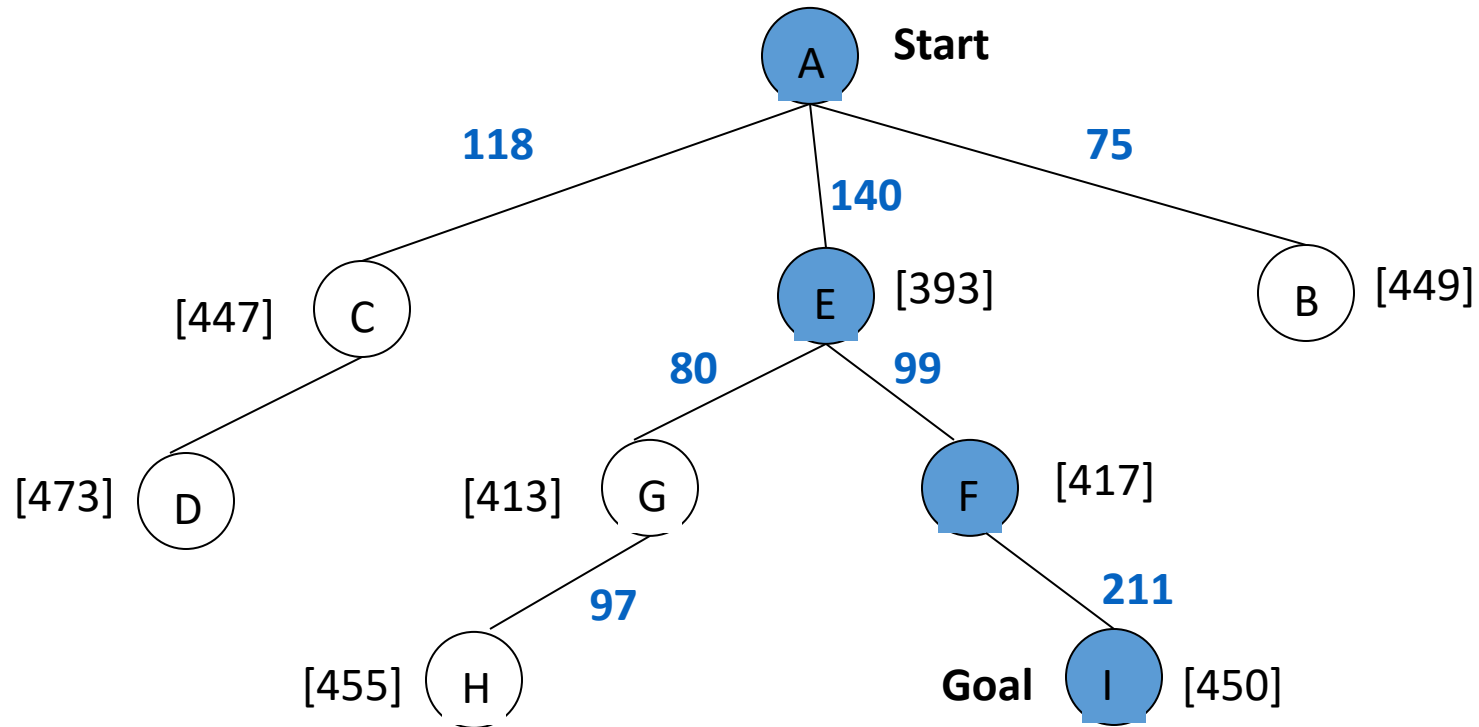
A* Search: Tree Search



A* Search: Tree Search



A* Search: Tree Search




A* not optimal !!!

$$\text{dist}(A-E-F-I) = 140 + 99 + 211 = 450$$


When heuristic h is not admissible

A* Tree Search


State space graph



Search tree



Consistency of Heuristics



- Main idea: Estimated heuristic costs \leq actual costs
 - Admissibility:
heuristic cost \leq actual cost to goal
 $h(A) \leq \text{actual cost from A to G}$
 - Consistency:
“heuristic step cost” \leq actual cost for each step
 $h(A) - h(C) \leq \text{cost}(A \text{ to } C)$
triangle inequality
 $h(A) \leq \text{cost}(A \text{ to } C) + h(C)$
- Consequences of consistency:
 - The f value along a path never decreases
 - A* graph search is optimal

But an **admissible heuristic is not always consistent**

A* Algorithm

1. Search queue Q is empty.
2. Place the start state s in Q with f value $h(s)$.
3. If Q is empty, return failure.
4. Take node n from Q with lowest f value.
(Keep Q sorted by f values and pick the first element).
5. If n is a goal node, stop and return solution.
6. Generate successors of node n.
7. For each successor n' of n do:
 - a) Compute $f(n') = g(n) + h(n')$.
 - b) If n' is new (never generated before), add n' to Q.
 - c) If node n' is already in Q with a higher f value, replace it with current $f(n')$ and place it in sorted order in Q.End for
8. Go back to step 3.

function UNIFORM-COST-SEARCH(**problem**) **returns** a solution, or failure

initialize the **explored set** to be empty

initialize the **frontier** as a priority queue using $g(n)$ as the priority

add initial state of **problem** to **frontier** with priority $g(S) = 0$

loop do

if the **frontier** is empty **then**

return failure

 choose a **node** and remove it from the **frontier**

if the **node** contains a goal state **then**

return the corresponding solution

 add the **node** state to the **explored set**

 for each resulting **child** from node

if the **child** state is not already in the **frontier** or **explored set** **then**

 add **child** to the **frontier**

else if the **child** is already in the **frontier** with higher $g(n)$ **then**

 replace that **frontier** node with **child**

function A-STAR-SEARCH(**problem**) **returns** a solution, or failure

initialize the **explored set** to be empty

initialize the **frontier** as a priority queue using $f(n) = g(n) + h(n)$ as the priority

add initial state of **problem** to **frontier** with priority $f(S) = 0 + h(S)$

loop do

if the **frontier** is empty **then**

return failure

 choose a **node** and remove it from the **frontier**

if the **node** contains a goal state **then**

return the corresponding solution

 add the **node** state to the **explored set**

 for each resulting **child** from node

if the **child** state is not already in the **frontier** or **explored set** **then**

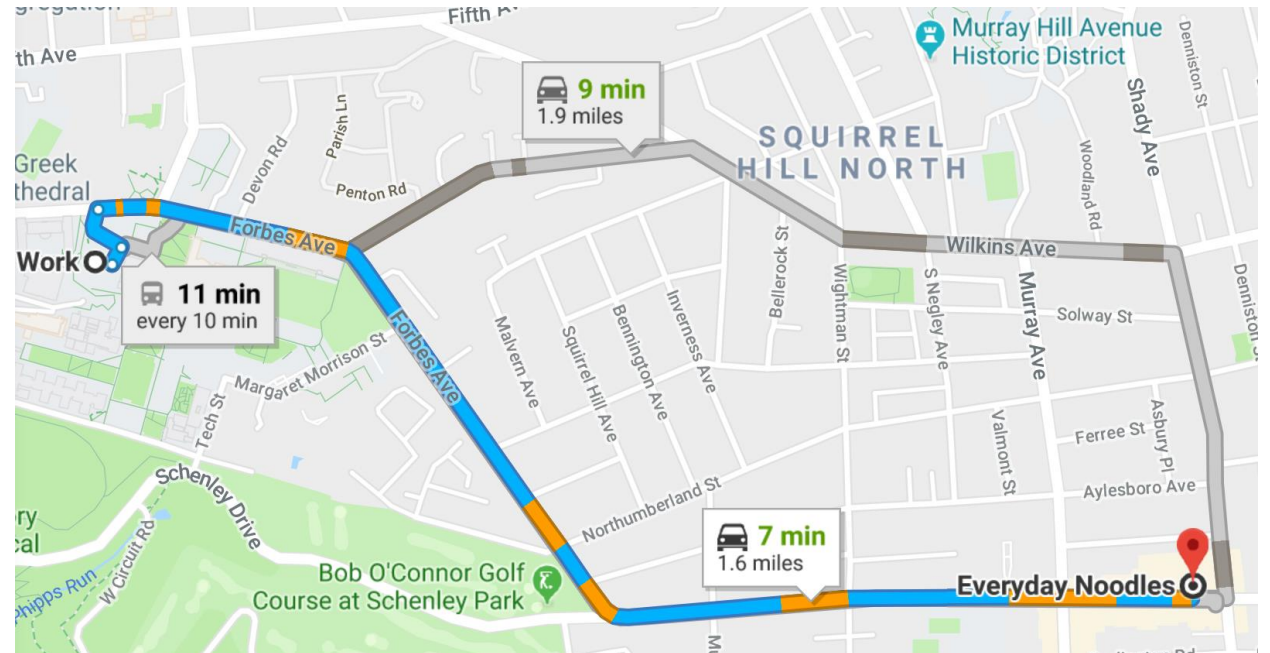
 add **child** to the **frontier**

else if the **child** is already in the **frontier** with higher $f(n)$ **then**

 replace that **frontier** node with **child**

A* Applications

- Pathing / routing problems
- Resource planning problems
- Robot motion planning
- Language analysis
- Video games
- Machine translation
- Speech recognition
- ...






Properties of A^*

Properties of A*


- Complete? **Yes**
 - (unless there are infinitely many nodes with $f \leq f(G)$)
- Time? **Exponential**
 - It depends on the **quality of heuristic** but is still exponential.
 - with respect to the length of the solution
 - better than other algorithms, but still problematic
- Space? Keeps all nodes in memory
 - **$O(b^d)$** space complexity,
 - A* keeps all generated nodes in memory
 - But an iterative deepening version is possible (IDA*).
- Optimal? **Yes**
 - A* is optimal if heuristic h is admissible (optimal).

Video of Demo Contours




Optimality

- Tree search:
 - A* is optimal if heuristic is admissible
 - UCS is a special case ($h = 0$)
- Graph search:
 - A* optimal if heuristic is consistent
 - UCS optimal ($h = 0$ is consistent)
- Consistency implies admissibility
- In general, most natural admissible heuristics tend to be consistent, especially if from relaxed problems



Optimality of A* Graph Search

- Sketch: consider what A* does with a consistent heuristic:
 - Fact 1: In tree search, A* expands nodes in increasing total f value (f-contours)
 - Fact 2: For every state s, nodes that reach s optimally are expanded before nodes that reach s suboptimally




Relationship of Search Algorithms

- Notation
 - $g(n)$ = known cost so far to reach n
 - $h(n)$ = estimated (optimal) cost from n to goal
 - $f(n) = g(n) + h(n)$ = estimated (optimal) total cost through n
- Uniform cost search: sort frontier by $g(n)$
- Greedy best-first search: sort frontier by $h(n)$
- A* search: sort frontier by $f(n)$
 - Optimal for admissible / consistent heuristics
 - Generally the preferred heuristic search framework
 - Memory-efficient versions of A* are available: RBFS, SMA*

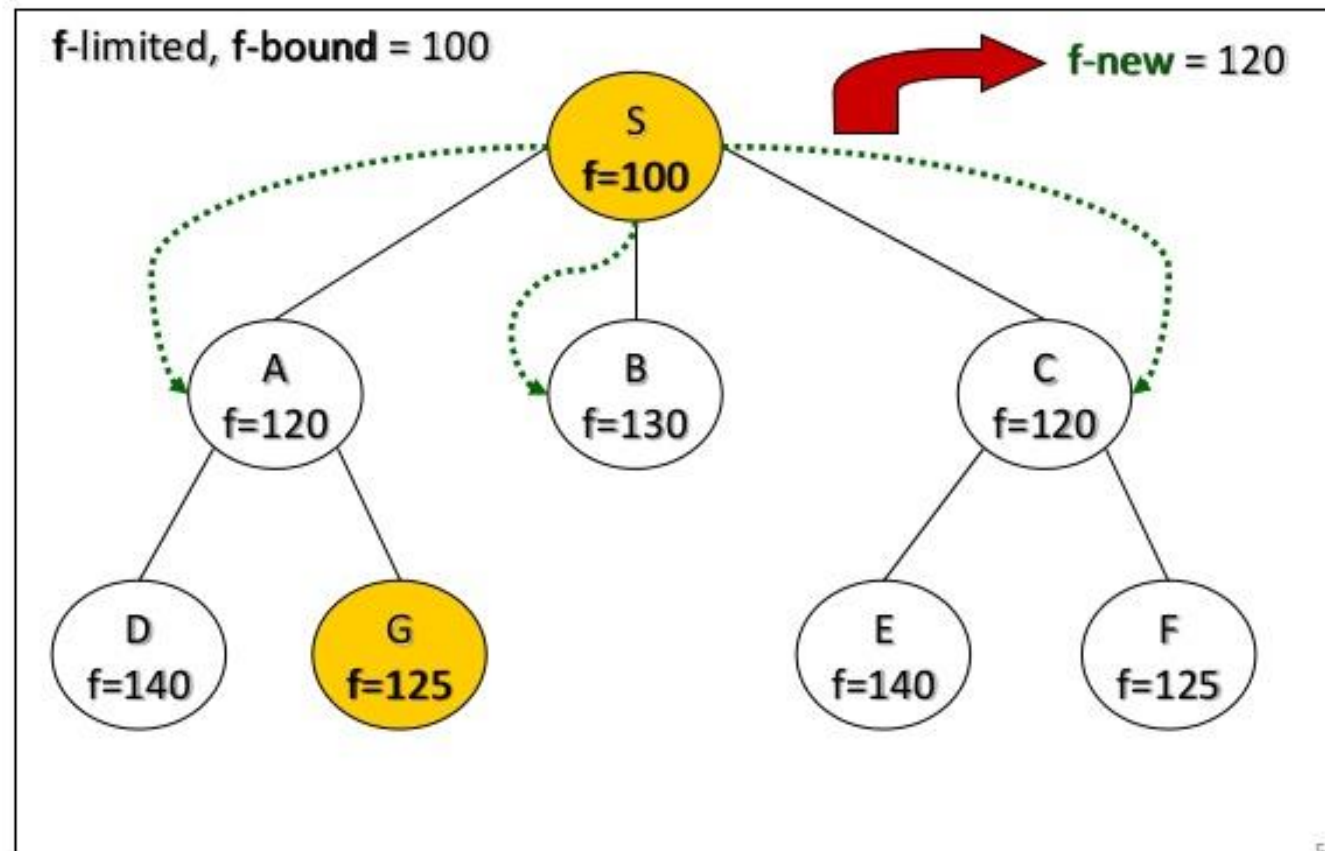
Question

- Perform Uniform Cost Search, the Best-First Search, and the A* algorithm. Here we suppose that **A is the initial node, and E is the target**. The cost of each edge and the heuristic value of the each node are also given in the figure.



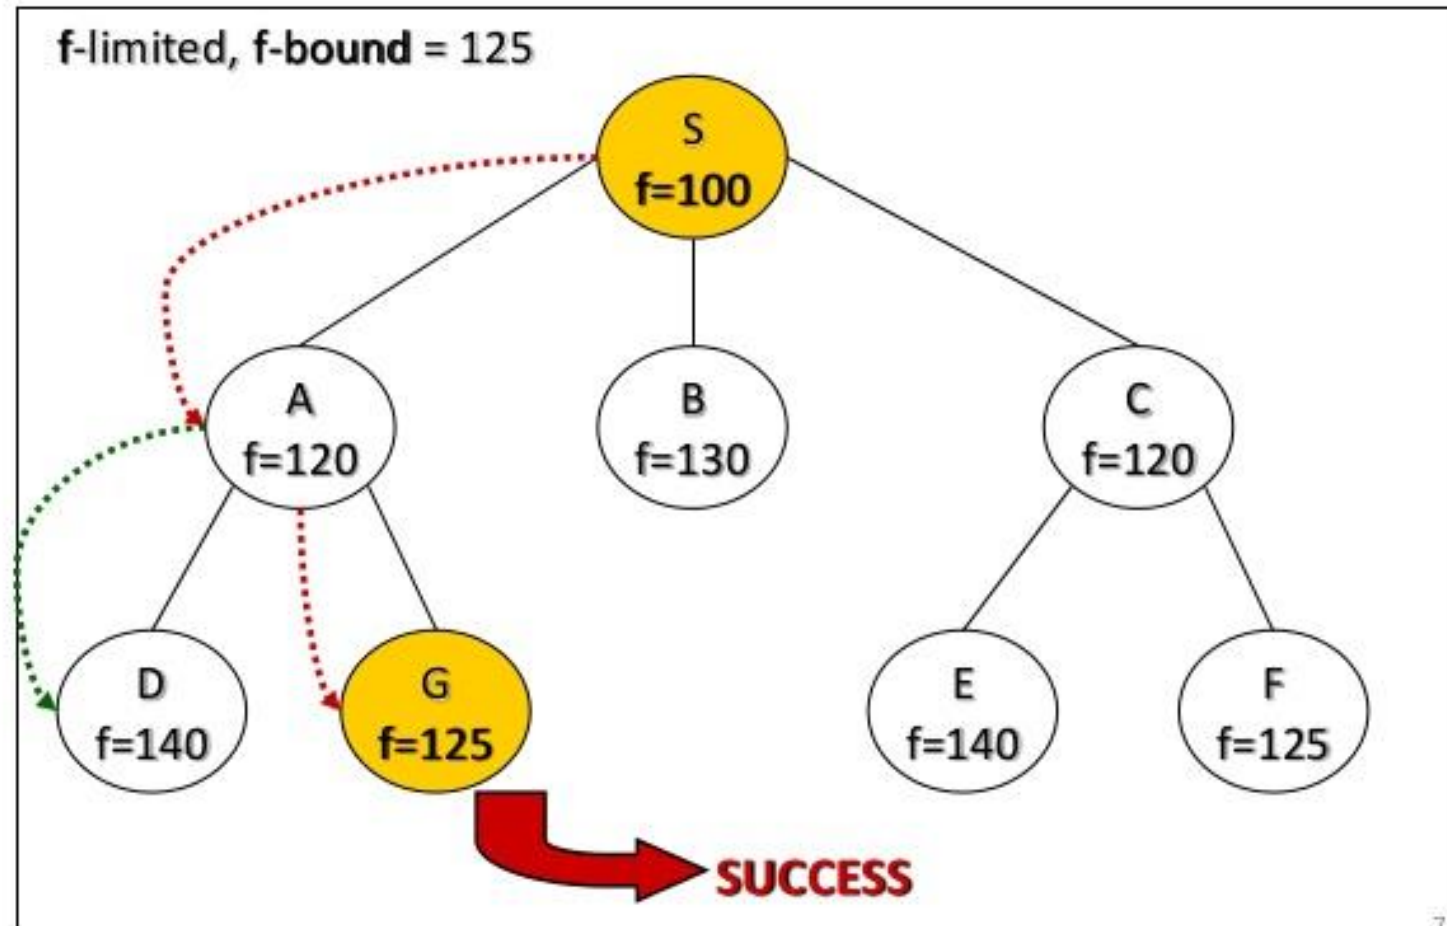
IDA*


Example





IDA*

Example









	S	A	B	C	D	G
heuristic	0	0	4	3	0	0




$f_{\text{bound}} = 0$
 $f_{\text{new}} = 7$



f_bound = 7
f_new = 10





f_bound = 11
f_new =

